

Appendix E: Experiment Run Logs, Cost/Latency Logs and Results Evidence

E.1 Purpose

The appendix records the final main-run evidence index, validation status, prompt hashes, telemetry, scoring outputs, classification metrics, confusion matrices and qualitative error evidence.

E.2 Evidence Provenance and Curated Result Baseline

Appendix E links to the final main-run evidence packaged in `Dissertation_Submission_Artefacts.zip` under `03_Main_Run_Evidence/`. The primary audit route is `03_Main_Run_Evidence/Appendix_E_Evidence_Pack/`, which contains the curated main-run evidence. All summary tables in this appendix are derived from `03_Main_Run_Evidence/Derived_Analysis/Main_Run_Results_Extraction_and_Analysis_Workbook_v1.xlsx`, which serves as the derived-analysis baseline and records the locked evidence packages in its `Source_Files` sheet.

For reproducibility, the locked source packages are retained in `03_Main_Run_Evidence/Locked_Source_Zips/`. Runtime outputs from `04_Codebase/CS958-ROUS04-Codebase.zip/CS958-ROUS04-HarjinderSingh/outputs/` are provided only as supplementary context; the Appendix E evidence pack and the workbook form the canonical audit path. Package-level SHA-256 checksums are recorded in the workbook's `Source_Files` sheet and in the submitted evidence manifests for integrity verification.

Table E.1: Evidence layers used

Evidence layer	Submitted location	Role in Appendix
Curated Appendix E evidence pack	<code>03_Main_Run_Evidence/Appendix_E_Evidence_Pack/</code>	Primary Appendix E evidence route containing selected final main-run evidence files exported for Appendix E review.
Derived analysis baseline	<code>03_Main_Run_Evidence/Derived_Analysis/Main_Run_Results_Extraction_and_Analysis_Workbook_v1.xlsx</code>	Immediate source for Appendix E summary tables; prepared after the main run from locked evidence packages without reruns or rescoreing.
Locked source packages	<code>03_Main_Run_Evidence/Locked_Source_Zips/</code>	Full-live, scoring and Appendix E export ZIPs retained for audit, reproducibility and source-package verification.
Codebase runtime-output context	<code>04_Codebase/CS958-ROUS04-Codebase.zip/CS958-ROUS04-HarjinderSingh/outputs/</code>	Runtime-output tree used for reproducibility checks; supplementary context, not the main Appendix E navigation route.

E.3 Locked Inputs and Main-Run Completion Checks

This section verifies that all locked experiment inputs were complete before scoring and that the main run executed successfully. The checks cover the model-call queue, scoring key, RAG-context records, cleaned-filing inputs, telemetry rows, completed calls, retry count and validation issues. All values are taken from the curated Appendix E evidence pack and the derived-analysis workbook.

Table E.2: Locked-input and main-run completion checks

Check	Recorded value	Evidence source
Validation status	PASS	Appendix_E_Evidence_Pack/validation_report.md
Model-call queue rows	96	Appendix_E_Evidence_Pack/validation_report.md
Scoring-key rows	48	Appendix_E_Evidence_Pack/validation_report.md
RAG context rows	48	Appendix_E_Evidence_Pack/validation_report.md
Clean text files	6	Appendix_E_Evidence_Pack/validation_report.md
Telemetry rows retained for analysis	96	Workbook Sheet: Dashboard, Telemetry
Long_Context context-delivery-condition calls	48	Workbook Sheet: Dashboard, Telemetry
RAG-based context-delivery-condition calls	48	Workbook Sheet: Dashboard, Telemetry
LIVE_CALL_COMPLETED rows	96	Workbook Sheet: Dashboard, Telemetry
Total retry attempts used	0	Workbook Sheet: Telemetry
Validation issues recorded	None	Appendix_E_Evidence_Pack/validation_report.md

E.4 Prompt, Payload and Output Artefact Evidence

Appendix B defines the logging and artefact-retention design, and Appendix D defines the prompt templates and JSON schema. This section identifies the evidence files used to verify that all 96 condition-level calls have corresponding prompt hashes, payload-size records, request payloads, raw model outputs and parsed structured outputs. Full model-facing payloads and raw outputs are retained in 03_Main_Run_Evidence/Locked_Source_Zips/, keeping Appendix E as an evidence index rather than a payload archive. Locked-input artefact hashes are recorded in the Source_Files sheet of Main_Run_Results_Extraction_and_Analysis_Workbook_v1.xlsx and in the submitted evidence manifests.

Table E.3: Prompt, payload and model-output evidence files

Evidence file	Evidence location	Rows	Use
---------------	-------------------	------	-----

Evidence file	Evidence location	Rows	Use
prompt_hashes.csv	03_Main_Run_Evidence/Appendix_E_Evidence_Pack/	96	Payload and user-prompt hashes for all context-delivery-condition calls.
prompt_size_summary.csv	03_Main_Run_Evidence/Appendix_E_Evidence_Pack/	96	Prompt and payload size measurements for all context-delivery-condition calls.
full_model_request_payloads.jsonl	03_Main_Run_Evidence/Locked_Source_Zips/main_run_full_live_outputs_locked_v1.zip	96	Complete request payload records retained in the locked full-live source package.
raw_model_outputs.jsonl	03_Main_Run_Evidence/Locked_Source_Zips/main_run_full_live_outputs_locked_v1.zip	96	Provider raw response records for all completed calls.
parsed_model_outputs.jsonl	03_Main_Run_Evidence/Locked_Source_Zips/main_run_full_live_outputs_locked_v1.zip	96	Parsed structured-output records for all completed calls.

E.5 Telemetry, Token, Cost and Latency Evidence

The full 96-row per-call telemetry file is retained as 03_Main_Run_Evidence/Appendix_E_Evidence_Pack/main_run_telemetry.csv and in the locked full-live source package. Appendix E reports aggregate telemetry by context-delivery condition; the retained CSV and derived workbook remain the sources for individual-call token counts, latency, retry count and telemetry-estimated model-call cost.

Cost values were calculated from logged token counts using the GPT-5.5 Standard API pricing constants configured in the runtime settings. They are reported as telemetry-estimated model-call costs, not as provider-billed per-call charges.

Table E.4: Telemetry-estimated token, cost and latency summary by context-delivery condition

Context-delivery condition	Claim-Case-Calls	Input tokens	Output tokens	Total tokens	Telemetry-estimated cost USD	Cumulative elapsed seconds	Mean latency seconds	Median latency seconds
Long_Context	48	8,582,196	25,257	8,607,453	\$43.668690	497.17	10.36	9.19
RAG-based	48	443,289	20,111	463,400	\$2.819775	299.26	6.23	5.34
Total	96	9,025,485	45,368	9,070,853	\$46.488465	796.43	8.30	7.11

E.6 Scoring and Three-Label Classification Evidence

Scored predictions were compared with the locked scoring key after execution. The primary quantitative score is exact label match for the three-label classification task: *Supported*, *Contradicted* or *Insufficient Evidence*. The definitions below are used to interpret the scoring tables in this section.

Table E.5: Metric definitions used in scoring tables

Metric / term	Meaning in this dissertation
Correct	Predicted label exactly matches the locked expected label for a context-delivery-condition call.
Total	Number of scored context-delivery-condition calls in the relevant Long_Context or RAG-based group.
Accuracy	Correct divided by total scored context-delivery-condition calls.
Macro F1	Unweighted mean of the three label-category F1 values.
Precision	For a label category, the share of predictions assigned to that label category that were correct.
Recall	For a label category, the share of expected cases in that label category correctly predicted.
F1	Harmonic mean of precision and recall for a label category.
Expected-label count	Number of claim-cases in the scoring key whose expected label is the listed label category.
False positive	For a label category, a prediction assigned to that label category when the locked expected label was different.
False negative	For a label category, a claim-case whose locked expected label was that label category but the model predicted a different label.

Table E.6 provides context-delivery-condition accuracy and macro F1 values derived from Main_Run_Results_Extraction_and_Analysis_Workbook_v1.xlsx, sheet Condition_Metrics.

Table E.6: Context-delivery-condition accuracy and macro F1

Context-delivery condition	Correct	Total	Accuracy	Macro F1
Long_Context	41	48	85.42%	84.68%
RAG-based	42	48	87.50%	87.28%

Table E.7 provides label-category precision, recall and F1 values derived from Main_Run_Results_Extraction_and_Analysis_Workbook_v1.xlsx, sheet Label_Metrics. The workbook column named support is reported here as Expected-label count to avoid confusion with the Supported label category.

Table E.7: Label-category precision, recall and F1 by context-delivery condition

Context-delivery condition	Label category	Precision	Recall	F1	Expected-label count
Long_Context	Supported	100.00%	100.00%	100.00%	16
Long_Context	Contradicted	69.57%	100.00%	82.05%	16
Long_Context	Insufficient Evidence	100.00%	56.25%	72.00%	16

Context-delivery condition	Label category	Precision	Recall	F1	Expected-label count
RAG	Supported	100.00%	93.75%	96.77%	16
RAG	Contradicted	76.19%	100.00%	86.49%	16
RAG	Insufficient Evidence	91.67%	68.75%	78.57%	16

Table E.8 provides confusion-matrix counts by context-delivery condition, derived from Main_Run_Results_Extraction_and_Analysis_Workbook_v1.xlsx, sheet Confusion_Matrix. The workbook stores these counts in long format; the table presents them in pivoted form for readability.

Table E.8: Confusion-matrix counts by context-delivery condition

Context-delivery condition	Expected label category	Predicted: Supported	Predicted: Contradicted	Predicted: Insufficient Evidence
Long_Context	Supported	16	0	0
Long_Context	Contradicted	0	16	0
Long_Context	Insufficient Evidence	0	7	9
RAG	Supported	15	0	1
RAG	Contradicted	0	16	0
RAG	Insufficient Evidence	0	5	11

Table E.9 provides false-positive and false-negative counts by label category, derived from Main_Run_Results_Extraction_and_Analysis_Workbook_v1.xlsx, sheet Label_Metrics, and cross-checked against sheet Confusion_Matrix.

Table E.9: False-positive and false-negative summary by label category

Context-delivery condition	Label category treated as positive	False positives	False negatives	Observed result
Long_Context	Supported	0	0	No Supported false positives or false negatives.
Long_Context	Contradicted	7	0	Seven Insufficient Evidence claim-cases were over-called as Contradicted; no true Contradicted claim-cases were missed.
Long_Context	Insufficient Evidence	0	7	Seven true Insufficient Evidence claim-cases were missed and predicted as Contradicted.
RAG	Supported	0	1	One Supported claim-case was missed and predicted as Insufficient Evidence.
RAG	Contradicted	5	0	Five Insufficient Evidence claim-cases were over-called as Contradicted; no true Contradicted claim-cases were missed.
RAG	Insufficient Evidence	1	5	Five true Insufficient Evidence claim-cases were missed; one Supported claim-case was incorrectly predicted as Insufficient Evidence.

E.7 Qualitative Error Evidence

Qualitative error evidence supports interpretation of the quantitative findings without changing the scoring rule. Misclassified claim-cases are identified from the scored-prediction evidence; error-family coding is used only as post-scoring diagnosis.

Table E.10 provides error counts by context-delivery condition, derived from Main_Run_Results_Extraction_and_Analysis_Workbook_v1.xlsx, sheet Condition_Metrics.

Table E.10: Error count by context-delivery condition

Context-delivery condition	Incorrect calls (misclassified-claims)	Total scored calls
Long_Context	7	48
RAG-based	6	48

Table E.11 provides error counts by expected label category and context-delivery condition, derived from Main_Run_Results_Extraction_and_Analysis_Workbook_v1.xlsx, sheets Label_Metrics and Error_Rows.

Table E.11: Error count by expected label category and context-delivery condition

Context-delivery condition	Expected label category	Incorrect calls
Long_Context	Supported	0
Long_Context	Contradicted	0
Long_Context	Insufficient Evidence	7
RAG-based	Supported	1
RAG-based	Contradicted	0
RAG-based	Insufficient Evidence	5

Table E.12 provides filing-level error counts by company, derived from Main_Run_Results_Extraction_and_Analysis_Workbook_v1.xlsx, sheet Filing_Performance, and aggregated across the two context-delivery conditions.

Table E.12: Error count by filing and company

Filing ID	Company	Long_Context errors	RAG-based errors	Total errors
F001	Apple Inc.	1	0	1
F002	Amazon.com, Inc.	1	2	3

Filing ID	Company	Long_Context errors	RAG-based errors	Total errors
F003	JPMorgan Chase & Co.	1	1	2
F004	Bank of America Corp.	0	0	0
F005	Verizon Communications Inc.	2	1	3
F006	Ford Motor Co.	2	2	4

Table E.12a records the locked claim-case type distribution used as the denominator for interpreting case-type error patterns. These counts describe the dataset design and are not model-performance results.

Table E.12a: Locked claim-case type distribution used for error-pattern interpretation

Claim-case type	Locked claim-case count	Interpretation use
cross_section_synthesis	6	Distributed evidence across filing sections.
numeric_consistency	12	Numeric or table-based consistency checks.
risk_to_mda_linkage	5	Links between risk disclosure and MD&A / related disclosure.
segment_consistency	6	Segment, business-line or reporting-structure consistency checks.
causal_or_explanatory_claim	2	Claims involving explanation or causal attribution.
unsupported_causal_claim	6	Claims where causal support is absent or insufficient.
missing_cross_section_support	5	Claims requiring cross-section support that is missing or incomplete.
unsupported_forward_looking_or_external_claim	5	Claims involving unsupported forward-looking or external assertions.
business_model_and_segments	1	Business-model / segment-structure claim.
Total	48	Full locked main-run claim-case set.

Table E.12b provides error counts by claim-case type and context-delivery condition, derived from Main_Run_Results_Extraction_and_Analysis_Workbook_v1.xlsx, sheet Case_Type_Performance.

Table E.12b: Error count by claim-case type and context-delivery condition

Claim-case type	Long_Context errors	RAG-based errors	Total errors
missing_cross_section_support	1	0	1
risk_to_mda_linkage	0	1	1

Claim-case type	Long_Context errors	RAG-based errors	Total errors
unsupported_causal_claim	3	3	6
unsupported_forward_looking_or_external_claim	3	2	5
Total	7	6	13

Table E.12c: Error Concentration by Claim-Case Type

Claim-case type group	Locked claim-cases	Possible strategy outputs	Misclassified outputs	Finding
unsupported_causal_claim	6	12	6	High concentration
unsupported_forward_looking_or_external_claim	5	10	5	High concentration
missing_cross_section_support	5	10	1	Isolated
risk_to_mda_linkage	5	10	1	Isolated
all other claim-case types	27	54	0	No errors observed
Total	48	96	13	—

Appendix E, Section E.7, Table E.12d summarises the shared and single-strategy misclassification distribution.

Table E.12d: Shared and Single-Strategy Misclassification Distribution

Error grouping	Claim-case ID(s)	Filing / company	Strategy affected	Misclassified outputs	Error category	Diagnostic meaning
Shared failure across both strategies	F002_C06	Amazon	Long_Context and RAG-based	2	Unsupported causal claim	Same unsupported-causal sufficiency failure occurred under both context-delivery strategies.
Shared failure across both strategies	F003_C08	JPMorgan Chase	Long_Context and RAG-based	2	Unsupported forward-looking or external claim	Same forward-looking sufficiency failure occurred under both strategies.
Shared failure across both strategies	F005_C06	Verizon	Long_Context and RAG-based	2	Unsupported causal claim	Same unsupported-causal sufficiency failure occurred under both strategies.
Shared failure across both strategies	F006_C07	Ford	Long_Context and RAG-based	2	Unsupported causal claim	Same unsupported-causal sufficiency failure occurred under both strategies.
Shared failure across both strategies	F006_C08	Ford	Long_Context and RAG-based	2	Unsupported forward-looking or external claim	Same forward-looking sufficiency failure occurred under both strategies.
Single-strategy failure	F001_C08	Apple	Long_Context only	1	Missing cross-section support	Full-filing access did not prevent a metric-specific sufficiency error.
Single-strategy failure	F005_C08	Verizon	Long_Context only	1	Unsupported forward-looking or external claim	Full-filing access did not prevent a forward-looking sufficiency error.
Single-strategy failure	F002_C03	Amazon	RAG-based only	1	Risk-to-MD&A linkage	Retrieved context omitted a required cross-section evidence component.

Error grouping	Claim-case ID(s)	Filing / company	Strategy affected	Misclassified outputs	Error category	Diagnostic meaning
Total shared failures	5 claim-cases	Multiple filings	Both strategies	10	Mixed	Shared errors dominated, indicating a common evidential-sufficiency boundary problem.
Total single-strategy failures	3 claim-cases	Multiple filings	One strategy only	3	Mixed	Single-strategy errors show strategy-specific failure routes.
Total misclassified outputs	8 distinct claim-cases	Multiple filings	Both strategies considered	13	Mixed	13 misclassified outputs arose from 8 distinct claim-cases.

Table E.13 summarises claim-case type error counts aggregated across both context-delivery conditions.

Table E.13: Error count by claim-case type

Claim-case type	Incorrect calls
missing_cross_section_support	1
risk_to_mda_linkage	1
unsupported_causal_claim	6
unsupported_forward_looking_or_external_claim	5

Table E.14 provides the qualitative error-family summary, derived from Main_Run_Results_Extraction_and_Analysis_Workbook_v1.xlsx, sheet Error_Families.

Table E.14: Error-family summary

Qualitative error family	Count	Share of errors
Insufficient Evidence over-called as Contradicted	12	92.31%
Supported evidence missed or underused	1	7.69%

Claim-case type and qualitative error family are distinct analytical fields. Claim-case type is inherited from the locked pre-execution manifest and describes the designed evidence-use demand of the claim-case. Qualitative error family is assigned after scoring and describes the observed failure in an incorrect output. Table E.15 joins these fields for each misclassified context-delivery-strategy call; it does not imply a deterministic mapping from claim-case type to qualitative error family.

Table E.15 lists the misclassified condition-level calls, derived from Main_Run_Results_Extraction_and_Analysis_Workbook_v1.xlsx, sheet Error_Rows.

Table E.15: Misclassified context-delivery-condition calls

Run ID	Claim-case ID	Context-delivery condition	Expected label	Predicted label	Claim-case type	Qualitative error family
MR008_LC	F001_C08	Long_Context	Insufficient Evidence	Contradicted	missing_cross_section_support	Insufficient Evidence over-called as Contradicted
MR011_RAG	F002_C03	RAG	Supported	Insufficient Evidence	risk_to_mda_linkage	Supported evidence missed or underused
MR014_LC	F002_C06	Long_Context	Insufficient Evidence	Contradicted	unsupported_causal_claim	Insufficient Evidence over-called as Contradicted
MR014_RAG	F002_C06	RAG	Insufficient Evidence	Contradicted	unsupported_causal_claim	Insufficient Evidence over-called as Contradicted
MR024_LC	F003_C08	Long_Context	Insufficient Evidence	Contradicted	unsupported_forward_looking_or_external_claim	Insufficient Evidence over-called as Contradicted
MR024_RAG	F003_C08	RAG	Insufficient Evidence	Contradicted	unsupported_forward_looking_or_external_claim	Insufficient Evidence over-called as Contradicted
MR038_LC	F005_C06	Long_Context	Insufficient Evidence	Contradicted	unsupported_causal_claim	Insufficient Evidence over-called as Contradicted
MR038_RAG	F005_C06	RAG	Insufficient Evidence	Contradicted	unsupported_causal_claim	Insufficient Evidence over-called as Contradicted
MR040_LC	F005_C08	Long_Context	Insufficient Evidence	Contradicted	unsupported_forward_looking_or_external_claim	Insufficient Evidence over-called as Contradicted
MR047_LC	F006_C07	Long_Context	Insufficient Evidence	Contradicted	unsupported_causal_claim	Insufficient Evidence over-called as Contradicted
MR047_RAG	F006_C07	RAG	Insufficient Evidence	Contradicted	unsupported_causal_claim	Insufficient Evidence over-called as Contradicted
MR048_LC	F006_C08	Long_Context	Insufficient Evidence	Contradicted	unsupported_forward_looking_or_external_claim	Insufficient Evidence over-called as Contradicted
MR048_RAG	F006_C08	RAG	Insufficient Evidence	Contradicted	unsupported_forward_looking_or_external_claim	Insufficient Evidence over-called as Contradicted

Table E.16 extends the row-level error record in Table E.15 by joining the locked claim-case manifest, scored predictions and frozen RAG contexts. It provides a concise qualitative diagnosis of each misclassified condition-level call without changing the locked expected labels, predictions or scores. All 13 misclassified outputs reported High confidence, so confidence did not identify these failures.

Table E.16: Case-Level Qualitative Error-Diagnosis Matrix

Run / case	Company / condition	Claim focus	Expected → predicted	Evidence state	Observed failure	SEC 10-K significance
MR008_LC / F001_C08	Apple / Long_Context	Exact iPhone units sold	Insufficient Evidence → Contradicted	Full filing contained net-sales values and product commentary, but no unit-sales count.	Absence of the requested metric was treated as direct conflict.	Adjacent financial measures are not interchangeable; the system must distinguish an undisclosed metric from a contradicted value.
MR011_RAG / F002_C03	Amazon / RAG	AWS AI services linked to MD&A AI uncertainty	Supported → Insufficient Evidence	Frozen context retrieved AWS/AI material but omitted the required MD&A passage on demand, expenses and benefits.	One side of a two-section evidence relationship was missing.	A cross-section claim can fail when retrieval supplies only one evidential component.
MR014_LC / F002_C06	Amazon / Long_Context	AI as primary cause of 12% sales growth	Insufficient Evidence → Contradicted	Full filing reported sales growth, AI investment and several disclosed sales drivers, but did not attribute consolidated growth primarily to AI.	Alternative disclosed drivers were treated as disproof of an unsupported causal attribution.	Related disclosures do not establish causal primacy; absence of causal support is not automatically contradiction.
MR014_RAG / F002_C06	Amazon / RAG	AI as primary cause of 12% sales growth	Insufficient Evidence → Contradicted	Retrieved context contained sales drivers and AI-related material but no explicit primary-cause statement.	Same causal-sufficiency boundary failure under retrieved context.	The error persisted across context-delivery strategies, indicating a shared reasoning boundary rather than a purely retrieval-specific defect.
MR024_LC / F003_C08	JPMorgan Chase / Long_Context	Security and Resiliency Initiative guarantees lower 2026 losses	Insufficient Evidence → Contradicted	Full filing mentioned the initiative and forward-looking uncertainty but no guarantee of reduced operational losses.	Cautionary language was treated as direct contradiction of unsupported guarantee wording.	Guarantees, forecasts and risk controls require careful separation of unsupported certainty from explicit conflict.
MR024_RAG / F003_C08	JPMorgan Chase / RAG	Security and Resiliency Initiative guarantees lower 2026 losses	Insufficient Evidence → Contradicted	Retrieved context emphasised cyber/resiliency risks and uncertainty but did not establish the asserted guarantee.	Same forward-looking certainty boundary failure under retrieved context.	The shared error shows that uncertainty language can prompt over-confident contradiction in both conditions.
MR038_LC / F005_C06	Verizon / Long_Context	AI automation as primary cause of Consumer revenue growth	Insufficient Evidence → Contradicted	Full filing disclosed Consumer revenue growth, specific commercial drivers, and separate AI/network discussion.	Alternative operational drivers were treated as disproof of an unsupported AI-causality claim.	Financial-driver attribution requires explicit linkage, not thematic proximity between technology and revenue disclosures.
MR038_RAG / F005_C06	Verizon / RAG	AI automation as primary cause of Consumer revenue growth	Insufficient Evidence → Contradicted	Retrieved context contained revenue drivers and network-technology material but no explicit primary-cause attribution to AI.	Same causal-sufficiency boundary failure under retrieved context.	The repeated error indicates a shared tendency to convert unsupported causal specificity into contradiction.
MR040_LC / F005_C08	Verizon / Long_Context	Frontier transaction guarantees higher 2026 revenue	Insufficient Evidence → Contradicted	Full filing discussed the transaction, expected benefits and forward-looking risks, but no guaranteed 2026 revenue increase.	Risk and cautionary language was treated as direct contradiction of unsupported guarantee wording.	Transaction benefits and forward-looking expectations should not be converted into guarantees or definitive negatives without explicit evidence.
MR047_LC / F006_C07	Ford / Long_Context	Tariffs rather than special items primarily caused 2025 net loss	Insufficient Evidence → Contradicted	Full filing discussed tariff effects and stated that the net-income decline was primarily driven by special items and lower segment EBIT.	Direct alternative-driver evidence triggered contradiction although the locked rubric treated the causal-primacy claim as insufficiently established.	This case exposes sensitivity at the boundary between unsupported causal primacy and evidence that appears to favour an alternative cause.
MR047_RAG / F006_C07	Ford / RAG	Tariffs rather than special items primarily caused 2025 net loss	Insufficient Evidence → Contradicted	Retrieved context included both tariff effects and special-item/segment explanations.	Same causal-boundary overcall under retrieved context.	The paired error indicates label-boundary difficulty rather than missing evidence alone.
MR048_LC / F006_C08	Ford / Long_Context	Model e guaranteed EBIT-positive in 2026	Insufficient Evidence → Contradicted	Full filing contained forward-looking guidance indicating a Model e EBIT loss and no guarantee.	Opposite outlook evidence was treated as contradiction, while the locked rubric classified the unsupported guarantee as Insufficient Evidence.	The case highlights sensitivity between unsupported certainty and explicit contrary outlook.

Run / case	Company / condition	Claim focus	Expected → predicted	Evidence state	Observed failure	SEC 10-K significance
MR048_RAG / F006_C08	Ford / RAG	Model e guaranteed EBIT-positive in 2026	Insufficient Evidence → Contradicted	Retrieved context included 2026 guidance indicating a Model e EBIT loss and forward-looking caveats.	Same guarantee-versus-outlook boundary overcall under retrieved context.	The shared error shows that strong contrary guidance can dominate the rubric's stricter treatment of unsupported guarantee language.

Source note: derived from the locked claim-case manifest, scoring key, scored_predictions.csv and rag_retrieved_contexts_locked_v1.jsonl. The matrix is interpretive audit evidence and does not alter the locked quantitative results.

E.8 Evidence Boundary

Appendix E reports main-run evidence only. Pilot validation, smoke-live checks, stress-smoke checks, runtime-patch evidence and earlier preparation-code material remain part of Appendix F or the audit/control artefacts. Appendix C remains the source of truth for claim-case text, expected labels, filing-reference anchors and ground-truth rationales. Appendix E records post-run evidence and analysis only; it does not modify the locked dataset or scoring key.